

lab-4

February 16, 2026

```
[53]: import tensorflow as tf
from tensorflow.keras.layers import Dense
import pandas as pd
```

```
[54]: data = pd.read_csv('data.csv')
data.head()
```

```
[54]: email      label
0  date wed NUMBER aug NUMBER NUMBER NUMBER NUMB...      0
1  martin a posted tassos papadopoulos the greek ...      0
2  man threatens explosion in moscow thursday aug...      0
3  klez the virus that won t die already the most...      0
4  in adding cream to spaghetti carbonara which ...      0
```

```
[55]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   email    2999 non-null   object 
 1   label    3000 non-null   int64  
dtypes: int64(1), object(1)
memory usage: 47.0+ KB
```

```
[56]: df = data.dropna(subset=['email'])
```

0.1 After removing 1 row

```
[57]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2999 entries, 0 to 2999
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   email    2999 non-null   object 

```

```
1    label    2999 non-null    int64
dtypes: int64(1), object(1)
memory usage: 70.3+ KB
```

0.2 Now applying TF-IDF vectorizer

```
[58]: from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer()
x = vectorizer.fit_transform(df['email'])
print(x.shape)
```

```
(2999, 34116)
```

0.3 Fitting Data in X and y

```
[59]: X_dense = x.toarray()
```

```
[60]: y = df['label'].to_numpy()
```

```
[61]: single_neuron = tf.keras.models.Sequential([
        tf.keras.layers.Dense(1, activation='sigmoid', input_shape=(x.shape[1],))
    ])
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[62]: single_neuron.compile(
        optimizer='adam',
        loss='binary_crossentropy',
        metrics=['accuracy']
    )
```

0.4 now apply Train test split

```
[63]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X_dense,
    y,
    test_size=0.2,
    random_state=42
)
```

0.5 Now train single neuron

```
[64]: history = single_neuron.fit(  
        X_train,  
        y_train,  
        epochs=10,  
        batch_size=32,  
        validation_split=0.2  
)
```

```
Epoch 1/10  
60/60          1s 9ms/step -  
accuracy: 0.7972 - loss: 0.6701 - val_accuracy: 0.8062 - val_loss: 0.6088  
Epoch 2/10  
60/60          1s 11ms/step -  
accuracy: 0.8433 - loss: 0.5827 - val_accuracy: 0.8062 - val_loss: 0.5454  
Epoch 3/10  
60/60          0s 6ms/step -  
accuracy: 0.8284 - loss: 0.5225 - val_accuracy: 0.8062 - val_loss: 0.4984  
Epoch 4/10  
60/60          0s 5ms/step -  
accuracy: 0.8319 - loss: 0.4725 - val_accuracy: 0.8062 - val_loss: 0.4615  
Epoch 5/10  
60/60          0s 5ms/step -  
accuracy: 0.8461 - loss: 0.4265 - val_accuracy: 0.8062 - val_loss: 0.4321  
Epoch 6/10  
60/60          0s 5ms/step -  
accuracy: 0.8502 - loss: 0.3980 - val_accuracy: 0.8125 - val_loss: 0.4075  
Epoch 7/10  
60/60          0s 5ms/step -  
accuracy: 0.8330 - loss: 0.3831 - val_accuracy: 0.8167 - val_loss: 0.3860  
Epoch 8/10  
60/60          0s 5ms/step -  
accuracy: 0.8624 - loss: 0.3480 - val_accuracy: 0.8208 - val_loss: 0.3672  
Epoch 9/10  
60/60          0s 6ms/step -  
accuracy: 0.8569 - loss: 0.3322 - val_accuracy: 0.8292 - val_loss: 0.3500  
Epoch 10/10  
60/60          0s 5ms/step -  
accuracy: 0.8814 - loss: 0.3078 - val_accuracy: 0.8354 - val_loss: 0.3346
```

```
[65]: import matplotlib.pyplot as plt  
  
plt.figure()  
  
# Train Loss (solid)  
plt.plot(history.history['loss'], linestyle='--')
```

```


# Validation Loss (solid)
plt.plot(history.history['val_loss'], linestyle='-' )

# Train Accuracy (dashed)
plt.plot(history.history['accuracy'], linestyle='--' )

# Validation Accuracy (dashed)
plt.plot(history.history['val_accuracy'], linestyle='---' )

plt.legend(['Train Loss', 'Val Loss', 'Train Acc', 'Val Acc'])
plt.xlabel("Epochs")
plt.ylabel("Value")
plt.title("Graph for LR = 0.001 (default)")
plt.show()

```



0.6 Graph on diffrent Learning Rate(LR)

```
[71]: import tensorflow as tf
import matplotlib.pyplot as plt

learning_rates = [0.1, 0.5, 0.01, 0.05, 0.001]

fig, axs = plt.subplots(1, 5, figsize=(20, 5))

for i, lr in enumerate(learning_rates):

    model = tf.keras.Sequential([
        tf.keras.layers.Dense(1, activation='sigmoid', input_shape=(X_train.
        ↪shape[1],))
    ])

    model.compile(
        optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
        loss='binary_crossentropy',
        metrics=['accuracy']
    )

    history = model.fit(
        X_train,
        y_train,
        epochs=10,
        batch_size=32,
        validation_split=0.2,
        verbose=0
    )


    axs[i].plot(history.history['loss'])
    axs[i].plot(history.history['val_loss'], linestyle='--')
    axs[i].plot(history.history['accuracy'])
    axs[i].plot(history.history['val_accuracy'], linestyle='--')

    axs[i].set_title(f"LR = {lr}")
    axs[i].set_xlabel("Epochs")
    axs[i].set_ylabel("Loss")

plt.tight_layout()
plt.legend(['Train Loss', 'Val Loss', 'Train Acc', 'Val Acc'])
plt.show()
```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```



```
[ ]:
```